

Aishwarya Iyer

AI Engineer — LLM Applications

Bengaluru, Karnataka (JP Nagar) | +91 98801 47362 | aishwarya.iyer.ai@gmail.com | linkedin.com/in/aishwaryaiyer-ai

SUMMARY

AI engineer with 6 years across ML and GenAI, the last two shipping retrieval-augmented LLM systems to production for an insurer and an NBFC. I care about the unglamorous parts: golden sets, regression gates, guardrails, and knowing why the answer changed before the user does. Notice period: 60 days.

SKILLS

GenAI: RAG design, structured output, prompt/context engineering, LLM-as-judge with human calibration, guardrails (PII redaction, groundedness checks), agent/tool-use patterns

Stack: Python, FastAPI, pgvector, OpenSearch, LangGraph, Azure OpenAI, vLLM (Llama/Mistral), MLflow

Engineering: pytest, Docker, Kubernetes, GitHub Actions, Terraform (working)

Compliance: DPDP Act-aligned PII handling, audit logging for model risk review

EXPERIENCE

AI Engineer — Cognivia Labs (GenAI consultancy), Bengaluru

Feb 2024 – present

- Built a claims-correspondence RAG assistant for a private life insurer, now handling 2,400 queries/day; groundedness gate blocks uncited answers.
- Designed the eval harness (240-item golden set, LLM-as-judge calibrated against human raters at kappa 0.81) — every prompt/model change ships through it in CI.
- Implemented PII redaction and audit logging to satisfy the client's DPO and model-risk team under the DPDP Act; system passed internal audit first pass.
- Cut inference cost 64% by routing simple intents to a fine-tuned small model behind the same API.
- Delivered a Hindi/Kannada intent-routing layer for an NBFC collections assistant (vernacular query normalization).

Machine Learning Engineer — Bazaario (e-commerce), Bengaluru

Sep 2021 – Jan 2024

- Owned the product-ranking model (LightGBM) serving 40M requests/day; +2.3% conversion in holdout.
- Built the team's first feature store and drift monitoring; two silent-degradation incidents caught in month one.
- Moved training pipelines to Kubernetes with MLflow tracking.

Decision Scientist — Mu Sigma, Bengaluru

Aug 2020 – Aug 2021

- Forecasting and churn models for a US retail client.

EDUCATION

M.Tech Computer Science & Engineering — IIIT Hyderabad, 2018 – 2020

B.E. Computer Science & Engineering — SSN College of Engineering, Chennai, 2014 – 2018

OPEN SOURCE

Maintainer of "granite-eval", a lightweight RAG evaluation toolkit (900+ GitHub stars)

LANGUAGES KNOWN

Tamil (native), English (fluent), Hindi (fluent), Kannada (working)